In [60]:
```python
#import all the necessary packages.

import PIL
import requests
from io import BytesIO
import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
import nltk
import math
import time
import re
import os
import seaborn as sns
from collections import Counter
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics import pairwise_distances
from matplotlib import gridspec
from scipy.sparse import hstack
import plotly
import plotly.figure_factory as ff
from plotly.graph_objs import Scatter, Layout

plotly.offline.init_notebook_mode(connected=True)
```

Text pre-processing

In [ ]:
```python
data = pd.read_pickle('pickels/16k_apperal_data')

# NLTK download stop words. [RUN ONLY ONCE]
# goto Terminal (Linux/Mac) or Command-Prompt (Window)
# In the temrinal, type these commands
# $python3
# $import nltk
# $nltk.download()
```

```
In [ ]:  # we use the list of stop words that are downloaded from nltk lib.
         stop_words = set(stopwords.words('english'))
         print ('list of stop words:', stop_words)

         def nlp_preprocessing(total_text, index, column):
             if type(total_text) is not int:
                 string = ""
                 for words in total_text.split():
                     # remove the special chars in review like '"#$@!%^&*()_+-~?>< etc.
                     word = ("".join(e for e in words if e.isalnum()))
                     # Conver all letters to lower-case
                     word = word.lower()
                     # stop-word removal
                     if not word in stop_words:
                         string += word + " "
                 data[column][index] = string
```

```
In [ ]:  start_time = time.clock()
         # we take each title and we text-preprocess it.
         for index, row in data.iterrows():
             nlp_preprocessing(row['title'], index, 'title')
         # we print the time it took to preprocess whole titles
         print(time.clock() - start_time, "seconds")
```

```
In [ ]:  data.head()
```

```
In [ ]:  data.to_pickle('pickels/16k_apperal_data_preprocessed')
```

Stemming

```
In [ ]:  from nltk.stem.porter import *
         stemmer = PorterStemmer()
         print(stemmer.stem('arguing'))
         print(stemmer.stem('fishing'))


         # We tried using stemming on our titles and it didnot work very well.
```

Text based product similarity

In [39]:
```python
data = pd.read_pickle('pickels/16k_apperal_data_preprocessed')
data.head()
```

Out[39]:

	asin	brand	color	medium_image_url	product_type_name	title	format
4	B004GSI2OS	FeatherLite	Onyx Black/ Stone	https://images-na.ssl-images-amazon.com/images...	SHIRT	featherlite ladies long sleeve stain resistant...	
6	B012YX2ZPI	HX-Kingdom Fashion T-shirts	White	https://images-na.ssl-images-amazon.com/images...	SHIRT	womens unique 100 cotton special olympics wor...	
15	B003BSRPB0	FeatherLite	White	https://images-na.ssl-images-amazon.com/images...	SHIRT	featherlite ladies moisture free mesh sport sh...	
27	B014ICEJ1Q	FNC7C	Purple	https://images-na.ssl-images-amazon.com/images...	SHIRT	supernatural chibis sam dean castiel neck tshi...	
46	B01NACPBG2	Fifth Degree	Black	https://images-na.ssl-images-amazon.com/images...	SHIRT	fifth degree womens gold foil graphic tees jun...	

In [40]:
```python
# Utility Functions which we will use through the rest of the workshop.
import PIL

#Display an image
def display_img(url,ax,fig):
    # we get the url of the apparel and download it
    response = requests.get(url)
    img = PIL.Image.open(BytesIO(response.content))
    # we will display it in notebook
    plt.imshow(img)

#plotting code to understand the algorithm's decision.
def plot_heatmap(keys, values, labels, url, text):
        # keys: list of words of recommended title
        # values: len(values) ==  len(keys), values(i) represents the occurence (
        # labels: len(labels) == len(keys), the values of labels depends on the n
                # if model == 'bag of words': labels(i) = values(i)
                # if model == 'tfidf weighted bag of words':labels(i) = tfidf(key
                # if model == 'idf weighted bag of words':labels(i) = idf(keys(i)
        # url : apparel's url

        # we will devide the whole figure into two parts
        gs = gridspec.GridSpec(2, 2, width_ratios=[4,1], height_ratios=[4,1])
        fig = plt.figure(figsize=(25,3))

        # 1st, ploting heat map that represents the count of commonly ocurred wor
        ax = plt.subplot(gs[0])
        # it displays a cell in white color if the word is intersection(lis of wo
        ax = sns.heatmap(np.array([values]), annot=np.array([labels]))
        ax.set_xticklabels(keys) # set that axis labels as the words of title
        ax.set_title(text) # apparel title

        # 2nd, plotting image of the the apparel
        ax = plt.subplot(gs[1])
        # we don't want any grid lines for image and no labels on x-axis and y-ax
        ax.grid(False)
        ax.set_xticks([])
        ax.set_yticks([])

        # we call dispaly_img based with paramete url
        display_img(url, ax, fig)

        # displays combine figure ( heat map and image together)
        plt.show()

def plot_heatmap_image(doc_id, vec1, vec2, url, text, model):

    # doc_id : index of the title1
    # vec1 : input apparels's vector, it is of a dict type {word:count}
    # vec2 : recommended apparels's vector, it is of a dict type {word:count}
    # url : apparels image url
    # text: title of recomonded apparel (used to keep title of image)
    # model, it can be any of the models,
        # 1. bag_of_words
        # 2. tfidf
        # 3. idf
```

```python
        # we find the common words in both titles, because these only words contribut
        intersection = set(vec1.keys()) & set(vec2.keys())

        # we set the values of non intersecting words to zero, this is just to show
        for i in vec2:
            if i not in intersection:
                vec2[i]=0

        # for labeling heatmap, keys contains list of all words in title2
        keys = list(vec2.keys())
        #  if ith word in intersection(lis of words of title1 and list of words of t
        values = [vec2[x] for x in vec2.keys()]

        # labels: len(labels) == len(keys), the values of labels depends on the model
            # if model == 'bag of words': labels(i) = values(i)
            # if model == 'tfidf weighted bag of words':labels(i) = tfidf(keys(i))
            # if model == 'idf weighted bag of words':labels(i) = idf(keys(i))

        if model == 'bag_of_words':
            labels = values
        elif model == 'tfidf':
            labels = []
            for x in vec2.keys():
                # tfidf_title_vectorizer.vocabulary_  it contains all the words in the
                # tfidf_title_features[doc_id, index_of_word_in_corpus] will give the
                if x in  tfidf_title_vectorizer.vocabulary_:
                    labels.append(tfidf_title_features[doc_id, tfidf_title_vectorizer
                else:
                    labels.append(0)
        elif model == 'idf':
            labels = []
            for x in vec2.keys():
                # idf_title_vectorizer.vocabulary_  it contains all the words in the
                # idf_title_features[doc_id, index_of_word_in_corpus] will give the
                if x in  idf_title_vectorizer.vocabulary_:
                    labels.append(idf_title_features[doc_id, idf_title_vectorizer.vo
                else:
                    labels.append(0)

        plot_heatmap(keys, values, labels, url, text)


# this function gets a list of wrods along with the frequency of each
# word given "text"
def text_to_vector(text):
    word = re.compile(r'\w+')
    words = word.findall(text)
    # words stores list of all words in given string, you can try 'words = text.s
    return Counter(words) # Counter counts the occurence of each word in list, it



def get_result(doc_id, content_a, content_b, url, model):
    text1 = content_a
    text2 = content_b
```

```
    # vector1 = dict{word11:#count, word12:#count, etc.}
    vector1 = text_to_vector(text1)

    # vector1 = dict{word21:#count, word22:#count, etc.}
    vector2 = text_to_vector(text2)

    plot_heatmap_image(doc_id, vector1, vector2, url, text2, model)
```

Bag of Words (BoW) on product titles.

In [41]:
```
from sklearn.feature_extraction.text import CountVectorizer
title_vectorizer = CountVectorizer()
title_features   = title_vectorizer.fit_transform(data['title'])
title_features.get_shape() # get number of rows and columns in feature matrix.
# title_features.shape = #data_points * #words_in_corpus
# CountVectorizer().fit_transform(corpus) returns
# the a sparse matrix of dimensions #data_points * #words_in_corpus

# title_features[doc_id, index_of_word_in_corpus] = number of times the word occu
```

Out[41]: (16042, 12609)

In [42]:
```python
def bag_of_words_model(doc_id, num_results):
    # doc_id: apparel's id in given corpus

    # pairwise_dist will store the distance from given input apparel to all rema
    # the metric we used here is cosine, the coside distance is mesured as K(X, 
    pairwise_dist = pairwise_distances(title_features,title_features[doc_id])

    # np.argsort will return indices of the smallest distances
    indices = np.argsort(pairwise_dist.flatten())[0:num_results]
    #pdists will store the smallest distances
    pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

    #data frame indices of the 9 smallest distace's
    df_indices = list(data.index[indices])

    for i in range(0,len(indices)):
        # we will pass 1. doc_id, 2. title1, 3. title2, url, model
        get_result(indices[i],data['title'].loc[df_indices[0]], data['title'].lo
        print('ASIN :',data['asin'].loc[df_indices[i]])
        print ('Brand:', data['brand'].loc[df_indices[i]])
        print ('Title:', data['title'].loc[df_indices[i]])
        print ('Euclidean similarity with the query image :', pdists[i])
        print('='*60)

#call the bag-of-words model for a product to get similar products.
bag_of_words_model(12566, 20) # change the index if you want to.
# In the output heat map each value represents the count value
# of the label word, the color represents the intersection
# with inputs title.

#try 12566
#try 931
```


burnt umber tiger tshirt zebra stripes xl xxl

```
ASIN : B00JXQB5FQ
Brand: Si Row
Title: burnt umber tiger tshirt zebra stripes xl  xxl
Euclidean similarity with the query image : 0.0
============================================================
```


pink tiger tshirt zebra stripes xl xxl

```
ASIN : B00JXQASS6
Brand: Si Row
Title: pink tiger tshirt zebra stripes xl  xxl
Euclidean similarity with the query image : 1.73205080757
```

TF-IDF based product similarity

```
In [43]: tfidf_title_vectorizer = TfidfVectorizer(min_df = 0)
          tfidf_title_features = tfidf_title_vectorizer.fit_transform(data['title'])
          # tfidf_title_features.shape = #data_points * #words_in_corpus
          # CountVectorizer().fit_transform(courpus) returns the a sparase matrix of dimen
          # tfidf_title_features[doc_id, index_of_word_in_corpus] = tfidf values of the wor
```

```
In [44]: def tfidf_model(doc_id, num_results):
             # doc_id: apparel's id in given corpus

             # pairwise_dist will store the distance from given input apparel to all rema
             # the metric we used here is cosine, the coside distance is mesured as K(X,
             pairwise_dist = pairwise_distances(tfidf_title_features,tfidf_title_features

             # np.argsort will return indices of 9 smallest distances
             indices = np.argsort(pairwise_dist.flatten())[0:num_results]
             #pdists will store the 9 smallest distances
             pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

             #data frame indices of the 9 smallest distace's
             df_indices = list(data.index[indices])

             for i in range(0,len(indices)):
                 # we will pass 1. doc_id, 2. title1, 3. title2, url, model
                 get_result(indices[i], data['title'].loc[df_indices[0]], data['title'].l
                 print('ASIN :',data['asin'].loc[df_indices[i]])
                 print('BRAND :',data['brand'].loc[df_indices[i]])
                 print ('Eucliden distance from the given image :', pdists[i])
                 print('='*125)
          tfidf_model(12566, 20)
          # in the output heat map each value represents the tfidf values of the label wor
```



```
ASIN : B00JXQB5FQ
BRAND : Si Row
Eucliden distance from the given image : 0.0
=============================================================================
=================================================
```



```
ASIN : B00JXQASS6
BRAND : Si Row
Eucliden distance from the given image : 0.753633191245
=============================================================================
```

IDF based product similarity

In [45]:
```python
idf_title_vectorizer = CountVectorizer()
idf_title_features = idf_title_vectorizer.fit_transform(data['title'])

# idf_title_features.shape = #data_points * #words_in_corpus
# CountVectorizer().fit_transform(courpus) returns the a sparase matrix of dimens
# idf_title_features[doc_id, index_of_word_in_corpus] = number of times the word
```

In [46]:
```python
def n_containing(word):
    # return the number of documents which had the given word
    return sum(1 for blob in data['title'] if word in blob.split())

def idf(word):
    # idf = log(#number of docs / #number of docs which had the given word)
    return math.log(data.shape[0] / (n_containing(word)))
```

In [47]:
```python
# we need to convert the values into float
idf_title_features  = idf_title_features.astype(np.float)

for i in idf_title_vectorizer.vocabulary_.keys():
    # for every word in whole corpus we will find its idf value
    idf_val = idf(i)

    # to calculate idf_title_features we need to replace the count values with th
    # idf_title_features[:, idf_title_vectorizer.vocabulary_[i]].nonzero()[0] wil
    for j in idf_title_features[:, idf_title_vectorizer.vocabulary_[i]].nonzero(

        # we replace the count values of word i in document j with  idf_value of
        # idf_title_features[doc_id, index_of_word_in_courpus] = idf value of wor
        idf_title_features[j,idf_title_vectorizer.vocabulary_[i]] = idf_val
```

```
In [48]: def idf_model(doc_id, num_results):
             # doc_id: apparel's id in given corpus

             # pairwise_dist will store the distance from given input apparel to all rema
             # the metric we used here is cosine, the coside distance is mesured as K(X, \
             # http://scikit-learn.org/stable/modules/metrics.html#cosine-similarity
             pairwise_dist = pairwise_distances(idf_title_features,idf_title_features[doc_

             # np.argsort will return indices of 9 smallest distances
             indices = np.argsort(pairwise_dist.flatten())[0:num_results]
             #pdists will store the 9 smallest distances
             pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

             #data frame indices of the 9 smallest distace's
             df_indices = list(data.index[indices])

             for i in range(0,len(indices)):
                 get_result(indices[i],data['title'].loc[df_indices[0]], data['title'].lo
                 print('ASIN :',data['asin'].loc[df_indices[i]])
                 print('Brand :',data['brand'].loc[df_indices[i]])
                 print ('euclidean distance from the given image :', pdists[i])
                 print('='*125)


         idf_model(12566,20)
         # in the output heat map each value represents the idf values of the label word,
```

burnt umber tiger tshirt zebra stripes xl xxl

7.7	9	6.5	2.2	6.4	5.9	2.7	3.7
burnt	umber	tiger	tshirt	zebra	stripes	xl	xxl

ASIN : B00JXQB5FQ
Brand : Si Row
euclidean distance from the given image : 0.0
===
===

pink tiger tshirt zebra stripes xl xxl

2.9	6.5	2.2	6.4	5.9	2.7	3.7
pink	tiger	tshirt	zebra	stripes	xl	xxl

ASIN : B00JXQASS6
Brand : Si Row
euclidean distance from the given image : 12.2050713112
===

Text Semantics based product similarity

In [49]:

```
# credits: https://www.kaggle.com/c/word2vec-nlp-tutorial#part-2-word-vectors
# Custom Word2Vec using your own text data.
# Do NOT RUN this code.
# It is meant as a reference to build your own Word2Vec when you have
# lots of data.

'''
# Set values for various parameters
num_features = 300     # Word vector dimensionality
min_word_count = 1     # Minimum word count
num_workers = 4        # Number of threads to run in parallel
context = 10           # Context window size
downsampling = 1e-3    # Downsample setting for frequent words

# Initialize and train the model (this will take some time)
from gensim.models import word2vec
print ("Training model...")
model = word2vec.Word2Vec(sen_corpus, workers=num_workers, \
            size=num_features, min_count = min_word_count, \
            window = context)

'''
```

Out[49]: '\n# Set values for various parameters\nnum_features = 300   # Word vector dim
ensionality   \nmin_word_count = 1   # Minimum word count
\nnum_workers = 4   # Number of threads to run in parallel\ncontext = 10
Context window size
\ndownsampling = 1e-3   # Downsample setting for frequent words\n\n# Initialize
and train the model (this will take some time)\nfrom gensim.models import word2
vec\nprint ("Training model...")\nmodel = word2vec.Word2Vec(sen_corpus, workers
=num_workers,   size=num_features, min_count = min_word_count,
window = context)\n   \n'

In [50]:

```
from gensim.models import Word2Vec
from gensim.models import KeyedVectors
import pickle

# in this project we are using a pretrained model by google
# its 3.3G file, once you load this into your memory
# it occupies ~9Gb, so please do this step only if you have >12G of ram
# we will provide a pickle file wich contains a dict ,
# and it contains all our courpus words as keys and  model[word] as values
# To use this code-snippet, download "GoogleNews-vectors-negative300.bin"
# from https://drive.google.com/file/d/0B7XkCwpI5KDYNlNUTTlSS21pQmM/edit
# it's 1.9GB in size.

'''
model = KeyedVectors.load_word2vec_format('GoogleNews-vectors-negative300.bin', b
'''

#if you do NOT have RAM >= 12GB, use the code below.
with open('word2vec_model', 'rb') as handle:
    model = pickle.load(handle)
```

In [51]:

```python
# Utility functions

def get_word_vec(sentence, doc_id, m_name):
    # sentence : title of the apparel
    # doc_id: document id in our corpus
    # m_name: model information it will take two values
        # if  m_name == 'avg', we will append the model[i], w2v representation o
        # if m_name == 'weighted', we will multiply each w2v[word] with the idf(
    vec = []
    for i in sentence.split():
        if i in vocab:
            if m_name == 'weighted' and i in  idf_title_vectorizer.vocabulary_:
                vec.append(idf_title_features[doc_id, idf_title_vectorizer.vocab
            elif m_name == 'avg':
                vec.append(model[i])
        else:
            # if the word in our courpus is not there in the google word2vec cor
            vec.append(np.zeros(shape=(300,)))
    # we will return a numpy array of shape (#number of words in title * 300 ) 3
    # each row represents the word2vec representation of each word (weighted/avg)
    return  np.array(vec)

def get_distance(vec1, vec2):
    # vec1 = np.array(#number_of_words_title1 * 300), each row is a vector of le
    # vec2 = np.array(#number_of_words_title2 * 300), each row is a vector of le

    final_dist = []
    # for each vector in vec1 we caluclate the distance(euclidean) to all vector
    for i in vec1:
        dist = []
        for j in vec2:
            # np.linalg.norm(i-j) will result the euclidean distance between vec
            dist.append(np.linalg.norm(i-j))
        final_dist.append(np.array(dist))
    # final_dist = np.array(#number of words in title1 * #number of words in titl
    # final_dist[i,j] = euclidean distance between vectors i, j
    return np.array(final_dist)


def heat_map_w2v(sentence1, sentence2, url, doc_id1, doc_id2, model):
    # sentance1 : title1, input apparel
    # sentance2 : title2, recommended apparel
    # url: apparel image url
    # doc_id1: document id of input apparel
    # doc_id2: document id of recommended apparel
    # model: it can have two values, 1. avg 2. weighted

    #s1_vec = np.array(#number_of_words_title1 * 300), each row is a vector(weigh
    s1_vec = get_word_vec(sentence1, doc_id1, model)
    #s2_vec = np.array(#number_of_words_title1 * 300), each row is a vector(weigh
    s2_vec = get_word_vec(sentence2, doc_id2, model)

    # s1_s2_dist = np.array(#number of words in title1 * #number of words in titl
    # s1_s2_dist[i,j] = euclidean distance between words i, j
    s1_s2_dist = get_distance(s1_vec, s2_vec)
```

```python
    # devide whole figure into 2 parts 1st part displays heatmap 2nd part display
    gs = gridspec.GridSpec(2, 2, width_ratios=[4,1],height_ratios=[2,1])
    fig = plt.figure(figsize=(15,15))

    ax = plt.subplot(gs[0])
    # ploting the heap map based on the pairwise distances
    ax = sns.heatmap(np.round(s1_s2_dist,4), annot=True)
    # set the x axis labels as recommended apparels title
    ax.set_xticklabels(sentence2.split())
    # set the y axis labels as input apparels title
    ax.set_yticklabels(sentence1.split())
    # set title as recommended apparels title
    ax.set_title(sentence2)

    ax = plt.subplot(gs[1])
    # we remove all grids and axis labels for image
    ax.grid(False)
    ax.set_xticks([])
    ax.set_yticks([])
    display_img(url, ax, fig)

    plt.show()
```

In [52]:
```python
# vocab = stores all the words that are there in google w2v model
# vocab = model.wv.vocab.keys() # if you are using Google word2Vec

vocab = model.keys()
# this function will add the vectors of each word and returns the avg vector of
def build_avg_vec(sentence, num_features, doc_id, m_name):
    # sentace: its title of the apparel
    # num_features: the lenght of word2vec vector, its values = 300
    # m_name: model information it will take two values
        # if  m_name == 'avg', we will append the model[i], w2v representation o
        # if m_name == 'weighted', we will multiply each w2v[word] with the idf(u

    featureVec = np.zeros((num_features,), dtype="float32")
    # we will intialize a vector of size 300 with all zeros
    # we add each word2vec(wordi) to this fetureVec
    nwords = 0

    for word in sentence.split():
        nwords += 1
        if word in vocab:
            if m_name == 'weighted' and word in  idf_title_vectorizer.vocabulary_
                featureVec = np.add(featureVec, idf_title_features[doc_id, idf_ti
            elif m_name == 'avg':
                featureVec = np.add(featureVec, model[word])
    if(nwords>0):
        featureVec = np.divide(featureVec, nwords)
    # returns the avg vector of given sentance, its of shape (1, 300)
    return featureVec
```

Average Word2Vec product similarity.

In [53]:
```python
doc_id = 0
w2v_title = []
# for every title we build a avg vector representation
for i in data['title']:
    w2v_title.append(build_avg_vec(i, 300, doc_id,'avg'))
    doc_id += 1

# w2v_title = np.array(# number of doc in courpus * 300), each row corresponds to
w2v_title = np.array(w2v_title)
```

In [54]:
```python
def avg_w2v_model(doc_id, num_results):
    # doc_id: apparel's id in given corpus

    # dist(x, y) = sqrt(dot(x, x) - 2 * dot(x, y) + dot(y, y))
    pairwise_dist = pairwise_distances(w2v_title, w2v_title[doc_id].reshape(1,-1

    # np.argsort will return indices of 9 smallest distances
    indices = np.argsort(pairwise_dist.flatten())[0:num_results]
    #pdists will store the 9 smallest distances
    pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

    #data frame indices of the 9 smallest distace's
    df_indices = list(data.index[indices])

    for i in range(0, len(indices)):
        heat_map_w2v(data['title'].loc[df_indices[0]],data['title'].loc[df_indic
        print('ASIN :',data['asin'].loc[df_indices[i]])
        print('BRAND :',data['brand'].loc[df_indices[i]])
        print ('euclidean distance from given input image :', pdists[i])
        print('='*125)


avg_w2v_model(12566, 20)
# in the give heat map, each cell contains the euclidean distance between words
```
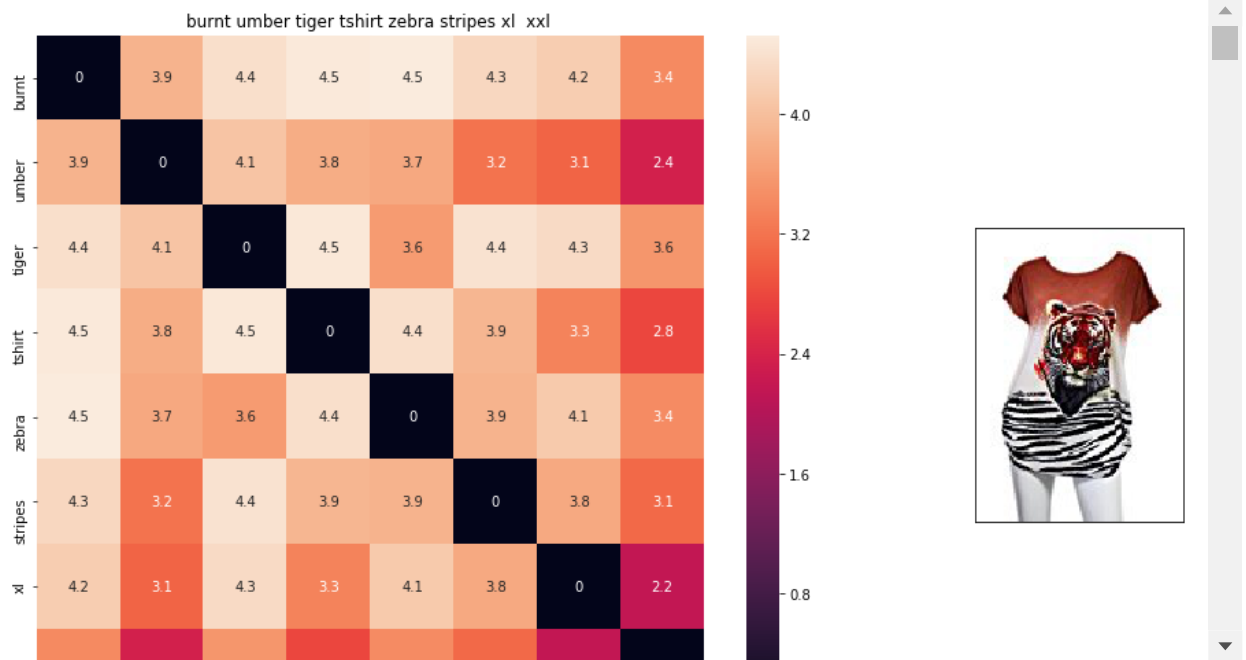
burnt umber tiger tshirt zebra stripes xl xxl

	burnt	umber	tiger	tshirt	zebra	stripes	xl	xxl
burnt	0	3.9	4.4	4.5	4.5	4.3	4.2	3.4
umber	3.9	0	4.1	3.8	3.7	3.2	3.1	2.4
tiger	4.4	4.1	0	4.5	3.6	4.4	4.3	3.6
tshirt	4.5	3.8	4.5	0	4.4	3.9	3.3	2.8
zebra	4.5	3.7	3.6	4.4	0	3.9	4.1	3.4
stripes	4.3	3.2	4.4	3.9	3.9	0	3.8	3.1
xl	4.2	3.1	4.3	3.3	4.1	3.8	0	2.2

IDF weighted Word2Vec for product similarity

In [55]:
```python
doc_id = 0
w2v_title_weight = []
# for every title we build a weighted vector representation
for i in data['title']:
    w2v_title_weight.append(build_avg_vec(i, 300, doc_id,'weighted'))
    doc_id += 1
# w2v_title = np.array(# number of doc in courpus * 300), each row corresponds to
w2v_title_weight = np.array(w2v_title_weight)
```

```
In [56]: def weighted_w2v_model(doc_id, num_results):
             # doc_id: apparel's id in given corpus

             # pairwise_dist will store the distance from given input apparel to all rema
             # the metric we used here cosine, the coside distance is mesured as K(X,
             # http://scikit-learn.org/stable/modules/metrics.html#cosine-similarity
             pairwise_dist = pairwise_distances(w2v_title_weight, w2v_title_weight[doc_id

             # np.argsort will return indices of 9 smallest distances
             indices = np.argsort(pairwise_dist.flatten())[0:num_results]
             #pdists will store the 9 smallest distances
             pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

             #data frame indices of the 9 smallest distace's
             df_indices = list(data.index[indices])

             for i in range(0, len(indices)):
                 heat_map_w2v(data['title'].loc[df_indices[0]],data['title'].loc[df_indic
                 print('ASIN :',data['asin'].loc[df_indices[i]])
                 print('Brand :',data['brand'].loc[df_indices[i]])
                 print('euclidean distance from input :', pdists[i])
                 print('='*125)

         weighted_w2v_model(12566, 20)
         #931
         #12566
         # in the give heat map, each cell contains the euclidean distance between words
```


Weighted similarity using brand and color.

In [57]:
```python
# some of the brand values are empty.
# Need to replace Null with string "NULL"
data['brand'].fillna(value="Not given", inplace=True )

# replace spaces with hypen
brands = [x.replace(" ", "-") for x in data['brand'].values]
types = [x.replace(" ", "-") for x in data['product_type_name'].values]
colors = [x.replace(" ", "-") for x in data['color'].values]

brand_vectorizer = CountVectorizer()
brand_features = brand_vectorizer.fit_transform(brands)

type_vectorizer = CountVectorizer()
type_features = type_vectorizer.fit_transform(types)

color_vectorizer = CountVectorizer()
color_features = color_vectorizer.fit_transform(colors)

extra_features = hstack((brand_features, type_features, color_features)).tocsr()
```

```
In [61]: def heat_map_w2v_brand(sentance1, sentance2, url, doc_id1, doc_id2, df_id1, df_i

             # sentance1 : title1, input apparel
             # sentance2 : title2, recommended apparel
             # url: apparel image url
             # doc_id1: document id of input apparel
             # doc_id2: document id of recommended apparel
             # df_id1: index of document1 in the data frame
             # df_id2: index of document2 in the data frame
             # model: it can have two values, 1. avg 2. weighted

             #s1_vec = np.array(#number_of_words_title1 * 300), each row is a vector(weigh
             s1_vec = get_word_vec(sentance1, doc_id1, model)
             #s2_vec = np.array(#number_of_words_title2 * 300), each row is a vector(weigh
             s2_vec = get_word_vec(sentance2, doc_id2, model)

             # s1_s2_dist = np.array(#number of words in title1 * #number of words in titl
             # s1_s2_dist[i,j] = euclidean distance between words i, j
             s1_s2_dist = get_distance(s1_vec, s2_vec)

             data_matrix = [['Asin','Brand', 'Color', 'Product type'],
                       [data['asin'].loc[df_id1],brands[doc_id1], colors[doc_id1], types
                       [data['asin'].loc[df_id2],brands[doc_id2], colors[doc_id2], types

             colorscale = [[0, '#1d004d'],[.5, '#f2e5ff'],[1, '#f2e5d1']] # to color the

             # we create a table with the data_matrix
             table = ff.create_table(data_matrix, index=True, colorscale=colorscale)
             # plot it with plotly
             plotly.offline.iplot(table, filename='simple_table')

             # devide whole figure space into 25 * 1:10 grids
             gs = gridspec.GridSpec(25, 15)
             fig = plt.figure(figsize=(25,5))

             # in first 25*10 grids we plot heatmap
             ax1 = plt.subplot(gs[:, :-5])
             # ploting the heap map based on the pairwise distances
             ax1 = sns.heatmap(np.round(s1_s2_dist,6), annot=True)
             # set the x axis labels as recommended apparels title
             ax1.set_xticklabels(sentance2.split())
             # set the y axis labels as input apparels title
             ax1.set_yticklabels(sentance1.split())
             # set title as recommended apparels title
             ax1.set_title(sentance2)

             # in last 25 * 10:15 grids we display image
             ax2 = plt.subplot(gs[:, 10:16])
             # we dont display grid lins and axis labels to images
             ax2.grid(False)
             ax2.set_xticks([])
             ax2.set_yticks([])

             # pass the url it display it
             display_img(url, ax2, fig)
```

```
plt.show()
```

In [62]:
```python
def idf_w2v_brand(doc_id, w1, w2, num_results):
    # doc_id: apparel's id in given corpus
    # w1: weight for  w2v features
    # w2: weight for brand and color features

    # pairwise_dist will store the distance from given input apparel to all rema
    # the metric we used here is cosine, the coside distance is mesured as K(X, `

    idf_w2v_dist  = pairwise_distances(w2v_title_weight, w2v_title_weight[doc_id
    ex_feat_dist = pairwise_distances(extra_features, extra_features[doc_id])
    pairwise_dist   = (w1 * idf_w2v_dist +  w2 * ex_feat_dist)/float(w1 + w2)

    # np.argsort will return indices of 9 smallest distances
    indices = np.argsort(pairwise_dist.flatten())[0:num_results]
    #pdists will store the 9 smallest distances
    pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

    #data frame indices of the 9 smallest distace's
    df_indices = list(data.index[indices])


    for i in range(0, len(indices)):
        heat_map_w2v_brand(data['title'].loc[df_indices[0]],data['title'].loc[df_
        print('ASIN :',data['asin'].loc[df_indices[i]])
        print('Brand :',data['brand'].loc[df_indices[i]])
        print('euclidean distance from input :', pdists[i])
        print('='*125)

idf_w2v_brand(12566, 5, 5, 20)
# in the give heat map, each cell contains the euclidean distance between words
```

In [63]:
```python
# brand and color weight =50
# title vector weight = 5

idf_w2v_brand(12566, 5, 50, 20)
```


ASIN : B00JXQB5FQ

Keras and Tensorflow to extract features

In [64]:
```python
import numpy as np
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dropout, Flatten, Dense
from keras import applications
from sklearn.metrics import pairwise_distances
import matplotlib.pyplot as plt
import requests
import PIL
import pandas as pd
import pickle
```

In [ ]:
```python
# https://gist.github.com/fchollet/f35fbc80e066a49d65f1688a7e99f069
# Code reference: https://blog.keras.io/building-powerful-image-classification-mo

# This code takes 40 minutes to run on a modern GPU (graphics card)
# like Nvidia  1050.
# GPU (NVidia 1050): 0.175 seconds per image

# This codse takes 160 minutes to run on a high end i7 CPU
# CPU (i7): 0.615 seconds per image.

#Do NOT run this code unless you want to wait a few hours for it to generate outp

# each image is converted into 25088 length dense-vector


# dimensions of our images.
img_width, img_height = 224, 224

top_model_weights_path = 'bottleneck_fc_model.h5'
train_data_dir = 'images2/'
nb_train_samples = 16042
epochs = 50
batch_size = 1


def save_bottlebeck_features():

    #Function to compute VGG-16 CNN for image feature extraction.

    asins = []
    datagen = ImageDataGenerator(rescale=1. / 255)

    # build the VGG16 network
    model = applications.VGG16(include_top=False, weights='imagenet')
    generator = datagen.flow_from_directory(
        train_data_dir,
        target_size=(img_width, img_height),
        batch_size=batch_size,
        class_mode=None,
        shuffle=False)

    for i in generator.filenames:
        asins.append(i[2:-5])

    bottleneck_features_train = model.predict_generator(generator, nb_train_sampl
    bottleneck_features_train = bottleneck_features_train.reshape((16042,25088))

    np.save(open('16k_data_cnn_features.npy', 'wb'), bottleneck_features_train)
    np.save(open('16k_data_cnn_feature_asins.npy', 'wb'), np.array(asins))


save_bottlebeck_features()
```

Visual features based product similarity.

In [66]:
```python
#load the features and corresponding ASINS info.
bottleneck_features_train = np.load('16k_data_cnn_features.npy')
asins = np.load('16k_data_cnn_feature_asins.npy')
asins = list(asins)

# load the original 16K dataset
data = pd.read_pickle('pickels/16k_apperal_data_preprocessed')
df_asins = list(data['asin'])


from IPython.display import display, Image, SVG, Math, YouTubeVideo


#get similar products using CNN features (VGG-16)
def get_similar_products_cnn(doc_id, num_results):
    doc_id = asins.index(df_asins[doc_id])
    pairwise_dist = pairwise_distances(bottleneck_features_train, bottleneck_fea

    indices = np.argsort(pairwise_dist.flatten())[0:num_results]
    pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

    for i in range(len(indices)):
        rows = data[['medium_image_url','title']].loc[data['asin']==asins[indices
        for indx, row in rows.iterrows():
            display(Image(url=row['medium_image_url'], embed=True))
            print('Product Title: ', row['title'])
            print('Euclidean Distance from input image:', pdists[i])
            print('Amazon Url: www.amzon.com/dp/'+ asins[indices[i]])

get_similar_products_cnn(12566, 20)
```



```
Product Title:  burnt umber tiger tshirt zebra stripes xl  xxl
Euclidean Distance from input image: 0.0625
Amazon Url: www.amzon.com/dp/B00JXQB5FQ
```



```
Product Title:  pink tiger tshirt zebra stripes xl  xxl
```

Euclidean Distance from input image: 30.0501
Amazon Url: www.amzon.com/dp/B00JXQASS6

Product Title: yellow tiger tshirt tiger stripes l
Euclidean Distance from input image: 41.2611
Amazon Url: www.amzon.com/dp/B00JXQCUIC

Product Title: brown white tiger tshirt tiger stripes xl xxl
Euclidean Distance from input image: 44.0002
Amazon Url: www.amzon.com/dp/B00JXQCWTO

Product Title: kawaii pastel tops tees pink flower design
Euclidean Distance from input image: 47.3825
Amazon Url: www.amzon.com/dp/B071FCWD97

Product Title: womens thin style tops tees pastel watermelon print
Euclidean Distance from input image: 47.7184
Amazon Url: www.amzon.com/dp/B01JUNHBRM

Product Title: kawaii pastel tops tees baby blue flower design
Euclidean Distance from input image: 47.9021
Amazon Url: www.amzon.com/dp/B071SBCY9W

Product Title: edv cheetah run purple multi xl
Euclidean Distance from input image: 48.0465
Amazon Url: www.amzon.com/dp/B01CUPYBM0

Product Title: danskin womens vneck loose performance tee xsmall pink ombre
Euclidean Distance from input image: 48.1019
Amazon Url: www.amzon.com/dp/B01F7PHXY8

Product Title: summer alpaca 3d pastel casual loose tops tee design
Euclidean Distance from input image: 48.1189
Amazon Url: www.amzon.com/dp/B01I80A93G

Product Title: miss chievous juniors striped peplum tank top medium shadowpeach
Euclidean Distance from input image: 48.1313
Amazon Url: www.amzon.com/dp/B0177DM70S

Product Title: red pink floral heel sleeveless shirt xl xxl
Euclidean Distance from input image: 48.1695
Amazon Url: www.amzon.com/dp/B00JV63QQE

Product Title: moana logo adults hot v neck shirt black xxl
Euclidean Distance from input image: 48.2568
Amazon Url: www.amzon.com/dp/B01LX6H43D

Product Title: abaday multicolor cartoon cat print short sleeve longline shirt large
Euclidean Distance from input image: 48.2656
Amazon Url: www.amzon.com/dp/B01CR57YY0

Product Title: kawaii cotton pastel tops tees peach pink cactus design
Euclidean Distance from input image: 48.3626
Amazon Url: www.amzon.com/dp/B071WYLBZS

Product Title: chicago chicago 18 shirt women pink
Euclidean Distance from input image: 48.3836
Amazon Url: www.amzon.com/dp/B01GXAZTRY

Product Title: yichun womens tiger printed summer tshirts tops
Euclidean Distance from input image: 48.4493
Amazon Url: www.amzon.com/dp/B010NN9RXO

Product Title: nancy lopez whimsy short sleeve whiteblacklemon drop xs
Euclidean Distance from input image: 48.4789
Amazon Url: www.amzon.com/dp/B01MPX6IDX

```
Product Title:  womens tops tees pastel peach ice cream cone print
Euclidean Distance from input image: 48.558
Amazon Url: www.amzon.com/dp/B0734GRKZL
```



```
Product Title:  uswomens mary j blige without tshirts shirt
Euclidean Distance from input image: 48.6144
Amazon Url: www.amzon.com/dp/B01M0XXFKK
```

Exercise:

A weighted euclidean distance model using Visual, Title, Brand and Color

We can try passing different weights into the w1(IDF-w2v), w2(Brand & Type), w3(Color), w4(Image) of idf_w2v_brand_color_img(doc_id, w1, w2, w3, w4, num_results)

Example: idf_w2v_brand_color_img(12566, 5, 4, 2, 1, 20) idf_w2v_brand_color_img(12566, 5, 5, 2, 1, 20) idf_w2v_brand_color_img(12566, 5, 4, 1, 2, 20) idf_w2v_brand_color_img(12566, 5, 5, 1, 1, 20) ... so on

In [67]:
```python
import numpy as np
from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Dropout, Flatten, Dense
from keras import applications
from sklearn.metrics import pairwise_distances
from sklearn.feature_extraction.text import CountVectorizer
import matplotlib.pyplot as plt
import requests
import pandas as pd
import pickle
```

In [68]:
```python
data = pd.read_pickle('pickels/16k_apperal_data_preprocessed')
#data.head()
```

In [69]:
```python
# some of the brand values are empty.
# Need to replace Null with string "NULL"
data['brand'].fillna(value="Not given", inplace=True )

# replace spaces with hypen
brands = [x.replace(" ", "-") for x in data['brand'].values]
types = [x.replace(" ", "-") for x in data['product_type_name'].values]
colors = [x.replace(" ", "-") for x in data['color'].values]

brand_vectorizer = CountVectorizer()
brand_features = brand_vectorizer.fit_transform(brands).tocsr()

type_vectorizer = CountVectorizer()
type_features = type_vectorizer.fit_transform(types).tocsr()

color_vectorizer = CountVectorizer()
color_features = color_vectorizer.fit_transform(colors).tocsr()
```

```python
from IPython.display import display, Image, SVG, Math, YouTubeVideo

def idf_w2v_brand_color_img(doc_id, w1, w2, w3, w4, num_results):
    # doc_id: apparel's id in given corpus
    # w1: weight for  w2v features
    # w2: weight for brand features
    # w3: weight for color features
    # w4: weight for image features

    # pairwise_dist will store the distance from given input apparel to all rema
    # the metric we used here is cosine, the coside distance is mesured as K(X,
    # http://scikit-learn.org/stable/modules/metrics.html#cosine-similarity
    #load the features and corresponding ASINS info.
    bottleneck_features_train = np.load('16k_data_cnn_features.npy')
    asins = np.load('16k_data_cnn_feature_asins.npy')
    asins = list(asins)

# load the original 16K dataset
    img_data = pd.read_pickle('pickels/16k_apperal_data_preprocessed')
    df_asins = list(img_data['asin'])

#   doc_id = asins.index(df_asins[doc_id])

    idf_w2v_dist  = pairwise_distances(w2v_title_weight, w2v_title_weight[doc_id
    brand_dist = pairwise_distances(brand_features, brand_features[doc_id])
    type_dist = pairwise_distances(type_features, type_features[doc_id])
    color_dist = pairwise_distances(color_features, color_features[doc_id])
    img_dist = pairwise_distances(bottleneck_features_train, bottleneck_features
    pairwise_dist   = (w1 * idf_w2v_dist +  w2 * (brand_dist + type_dist) + w3 *

    # np.argsort will return indices of 9 smallest distances
    indices = np.argsort(pairwise_dist.flatten())[0:num_results]
    #pdists will store the 9 smallest distances
    pdists  = np.sort(pairwise_dist.flatten())[0:num_results]

    #data frame indices of the 9 smallest distace's
    df_indices = list(data.index[indices])


    for i in range(0, len(indices)):
        heat_map_w2v_brand(data['title'].loc[df_indices[0]],data['title'].loc[df_
        print('ASIN :',data['asin'].loc[df_indices[i]])
        print('Brand :',data['brand'].loc[df_indices[i]])
        print('euclidean distance from input :', pdists[i])
        print('='*125)

idf_w2v_brand_color_img(12566, 5, 4, 2, 1, 20)
# in the give heat map, each cell contains the euclidean distance between words
```

burnt umber tiger tshirt zebra stripes xl xxl

ASIN : B00JXQB5FQ